

# CYRIL NWACHUKWU

## Senior AI Developer & Client Solutions Architect

Email: K1ngceework97@gmail.com | Phone: +234 814 948 7720  
Location: Lagos, Nigeria (Global Remote Ready)  
LinkedIn: [linkedin.com/in/cyril-nwachukwu-96b029173](https://www.linkedin.com/in/cyril-nwachukwu-96b029173) |  
GitHub: [github.com/Thebeardedambivert](https://github.com/Thebeardedambivert)  
Portfolio: [cyril-nwachukwu.vercel.app](https://cyril-nwachukwu.vercel.app)

### EXECUTIVE SUMMARY

**Senior AI Developer and Solutions Architect** with 4+ years of experience acting as the primary technical bridge between international business stakeholders and high-velocity engineering teams. Specialized in translating ambiguous client requirements into production-grade multi-agent architectures (**LangGraph v1**), **Model Context Protocol (MCP)** integrations, deterministic RAG pipelines, and scalable **FastAPI** services. Proven track record leading technical discovery workshops, conducting rigorous architecture reviews, and operationalizing cost-metered, observable AI systems with OpenTelemetry across E-commerce, PropTech, and Enterprise Workflow Automation.

### CORE TECHNICAL & LEADERSHIP COMPETENCIES

- AI & Multi-Agent Frameworks:** LangGraph v1, Model Context Protocol (MCP), Cyclic State Graphs, Deterministic RAG, Human-in-the-Loop, LLM Evaluation Harnesses, Prompt Engineering, Semantic Search (pgvector, Pinecone).
- Backend & Microservices:** Python (AsyncIO), FastAPI, Django, PostgreSQL, Redis, SQL Analytics, RESTful APIs, Docker, Google Cloud Run, CI/CD (GitHub Actions).
- Observability & Reliability:** OpenTelemetry, Span-level LLM Cost & Latency Attribution, Structured Logging, HMAC-SHA256 Security Gates, TDD & Pytest (200+ suite benchmarks).
- Client & Technical Leadership:** Client Discovery Calls, Requirement Scoping & PRDs, Architecture Blueprints, Cross-Functional Alignment, Code Reviews, Team Mentorship.

### FEATURED PRODUCTION AI PROJECTS

- Shopify AI Ops Agent - Autonomous E-Commerce Operations**

Python, LangGraph v1, 5 MCP Servers, SQL Analytics

  - Architected an autonomous, human-governed ecommerce operations agent built on LangGraph v1 and 5 Model Context Protocol (MCP) servers for live store operations.
  - Implemented deterministic SQL analytics and HMAC-SHA256 safety approval gates, eliminating LLM math hallucinations and preventing unauthorized database mutations.
- LangGraph Multi-Judge Report Agent & Real-Time Telemetry**

LangGraph, OpenTelemetry, Parallel Orchestration

  - Built an enterprise report generation agent featuring parallel reviewer nodes, dynamic model routing, and evaluated output gates for high-accuracy document synthesis.
  - Engineered a real-time OpenTelemetry streaming telemetry pipeline that broadcasts span-level node status and token generation events without blocking graph state transitions.
- AgentPipe - Autonomous Coding Pipeline (Open Source)**

LangGraph, PostgreSQL, OpenTelemetry, Pytest

  - Developed an open-source coding agent that consumes issue tickets, navigates repositories, generates patches, and proves correctness via automated test suites before opening Pull Requests.
  - Built a metered LLM client with OpenTelemetry cost/latency attribution, content-hash caching, a 3-state acceptance gate, and an LLM Judge evaluation harness backed by 237 CI tests.
- ApexOps Orchestrator - Distributed AI Lead Platform**

Python, Postgres Outbox, Event Sourcing, Evals

  - Engineered an event-driven lead orchestration system featuring signed webhooks, Postgres transactional outbox pattern, idempotent billing, and evaluated AI safety guardrails.
- CoachOS V5 - Enterprise AI Practice Management Portal**

n8n, Supabase, Gemini, React, Google Cloud Run

  - Deployed full-stack client operations platform automating contract delivery, payment webhooks, and CRM sync, cutting post-session manual admin to zero with automated AI intelligence reports.

PROFESSIONAL EXPERIENCE

Senior AI Developer & Solutions Consultant

June 2023 - Present

Independent Practice (International Startups & Clients)

Global Remote

- **Client Discovery & Technical Scoping:** Serve as primary technical lead for international founders; run discovery workshops, translate business objectives into detailed engineering PRDs, and present architecture options with clear trade-offs across cost, latency, and reliability.
- **Multi-Agent & RAG Architecture:** Architect cyclic multi-agent graphs (LangGraph v1) and deterministic RAG pipelines (Pinecone, Supabase pgvector) incorporating hybrid search, reranking, and human-in-the-loop gates for high-stakes workflows.
- **Observability & Cost Governance:** Implement production-wide OpenTelemetry tracing to monitor span-level LLM latency, token expenditure, and error boundaries, giving stakeholders transparent operational metrics.
- **Engineering Quality & Mentorship:** Conduct weekly architecture reviews, establish TDD standards with pytest, and mentor developers on clean modular API patterns in FastAPI.
- **AI Guardrails & Safety:** Designed evaluation harnesses and deterministic verification layers preventing model drift, prompt injections, and invalid state mutations.

Software Developer (Backend & Systems)

February 2021 - May 2023

UB4US Smart Solutions

Lagos, Nigeria

- **Team Leadership:** Led a team of 4 engineers through full SDLC iterations, establishing CI/CD automation and strict code quality standards.
- **High-Performance REST APIs:** Designed and deployed scalable backend services in Django REST Framework and Flask microservices, optimizing database queries and caching to reduce backend latency by 40%.
- **Workflow Automation:** Engineered custom internal integration pipelines between Google Workspace, Slack, and third-party APIs, saving 10+ operational hours per week.
- **Client-Facing Technical Delivery:** Managed client communication during technical onboarding, integration reviews, and feature delivery.

SYSTEM ARCHITECTURE & ENGINEERING INVARIANTS

- **Deep Modules & Information Hiding:** Enforce narrow, simple public interfaces for complex subgraphs, concealing internal node mechanics and state transitions.
- **Transactional Outbox & Event Sourcing:** Implement resilient event emission in distributed agent pipelines to ensure zero message loss and idempotent billing.
- **Model Routing & Cost Optimization:** Leverage per-role LLM routing (fast reasoning models for classification, heavy models for synthesis) with semantic caching to cut API spend by up to 60%.

EDUCATION

Bachelor of Arts (B.A.) in Philosophy - Obafemi Awolowo University

Graduated 2020

Specialization in Formal Logic, Epistemology, and First-Principles Problem Formulation.